

Shea-FA691-HW3

February 15, 2023

1 FA691 Homework 3

2 Due: Wednesday, February 15 @ 11:59PM

Name: Ryan Shea

Date: 2023-02-02

```
[ ]: import numpy as np
import yfinance as yf
import matplotlib.pyplot as plt
import pandas as pd

# Set seed of random number generator
CWID = 10445281 #Place here your Campus wide ID number, this will personalize
#your results, but still maintain the reproduceable nature of using seeds.
#If you ever need to reset the seed in this assignment, use this as your seed
#Papers that use -1 as this CWID variable will earn 0's so make sure you change
#this value before you submit your work.
personal = CWID % 10000
np.random.seed(personal)
```

2.1 Question 1 (10pt)

2.1.1 Question 1.1

Use the `yfinance` package (or other method of your choice) to obtain the daily adjusted close prices for SPY and IEF. You should have at least 5 years of data for both assets. Do **not** include any data after January 1, 2022. You should inspect the dates for your data to make sure you are including everything appropriately. Create a binary variable whether the SPY returns are above the IEF returns on a each day. Create a data frame (or array) of the daily log returns both both stocks along with the lagged returns (at least 2 lags) and your binary class variable. Use the `print` command to display your data.

```
[ ]: import datetime
tickers = ['SPY', 'IEF']
stop = datetime.date(2023, 2, 8)
start = stop - datetime.timedelta(365*5+40) # 5 years, 40 days for buffer
```

```

data = yf.download(tickers=tickers, start=start, end=stop)['Adj Close']

df = data.apply(lambda x: np.log(x / x.shift(1))) # log(t / t-1)

df['SPY_ABOVE_IEF'] = (df['SPY'] > df['IEF']).astype(int)

lags = 10

for t in tickers:
    for i in range(1, lags+1):
        df[f"{t}_{i}"] = df[t].shift(i)

df = df.dropna()

print(df.head(), end='\n\n')
print(f"{df.shape = }")

```

```

[*****100%*****] 2 of 2 completed

```

	IEF	SPY	SPY_ABOVE_IEF	SPY_1	SPY_2	SPY_3	\
Date							
2018-01-18	-0.003166	-0.001682	1	0.009486	-0.003424	0.006498	
2018-01-19	-0.002598	0.004540	1	-0.001682	0.009486	-0.003424	
2018-01-22	-0.000386	0.008098	1	0.004540	-0.001682	0.009486	
2018-01-23	0.002887	0.002120	0	0.008098	0.004540	-0.001682	
2018-01-24	-0.001924	-0.000388	1	0.002120	0.008098	0.004540	

	SPY_4	SPY_5	SPY_6	SPY_7	...	IEF_1	IEF_2	\
Date					...			
2018-01-18	0.007270	-0.001531	0.002261	0.001827	...	-0.002201	0.000478	
2018-01-19	0.006498	0.007270	-0.001531	0.002261	...	-0.003166	-0.002201	
2018-01-22	-0.003424	0.006498	0.007270	-0.001531	...	-0.002598	-0.003166	
2018-01-23	0.009486	-0.003424	0.006498	0.007270	...	-0.000386	-0.002598	
2018-01-24	-0.001682	0.009486	-0.003424	0.006498	...	0.002887	-0.000386	

	IEF_3	IEF_4	IEF_5	IEF_6	IEF_7	IEF_8	\
Date							
2018-01-18	-0.000670	0.000670	-0.000287	-0.004769	-0.000476	-0.001236	
2018-01-19	0.000478	-0.000670	0.000670	-0.000287	-0.004769	-0.000476	
2018-01-22	-0.002201	0.000478	-0.000670	0.000670	-0.000287	-0.004769	
2018-01-23	-0.003166	-0.002201	0.000478	-0.000670	0.000670	-0.000287	
2018-01-24	-0.002598	-0.003166	-0.002201	0.000478	-0.000670	0.000670	

	IEF_9	IEF_10
Date		
2018-01-18	-0.000475	0.001045

```

2018-01-19 -0.001236 -0.000475
2018-01-22 -0.000476 -0.001236
2018-01-23 -0.004769 -0.000476
2018-01-24 -0.000287 -0.004769

```

```
[5 rows x 23 columns]
```

```
df.shape = (1273, 23)
```

2.1.2 Question 1.2

Split your data into training and testing sets (80% training and 20% test). This split should be done so that the causal relationship is kept consistent (i.e., split data at a specific time).

Run a logistic regression of the binary variable (of SPY returns greater than IEF returns) as a function of the lagged returns (at least 2 lags) for both stocks. This should be of the form (assuming 2 lags) of $p_t = [1 + \exp(-[\beta_0 + \beta_{SPY,1}r_{SPY,t-1} + \beta_{SPY,2}r_{SPY,t-2} + \beta_{IEF,1}r_{IEF,t-1} + \beta_{IEF,2}r_{IEF,t-2}])]^{-1}$. Evaluate the performance of this model by printing the confusion matrix and accuracy on the test data.

```

[ ]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import confusion_matrix, accuracy_score

temp = df.drop(["IEF", "SPY"], axis=1)

X_train, X_test, y_train, y_test = train_test_split(temp.drop("SPY_ABOVE_IEF", axis=1), temp["SPY_ABOVE_IEF"], test_size=0.2)

lr = LogisticRegression()
lr.fit(X_train, y_train)
l_pred = lr.predict(X_test)

print(f"Intercept: {float(lr.intercept_)}", end='\n\n')

for coef, name in zip(lr.coef_[0], X_train.columns):
    print(f" \u03B2 {name:>6}: {coef}")

print()

print(confusion_matrix(y_test, l_pred))
print(f"Score: {accuracy_score(y_test, l_pred)}")

```

```
Intercept: 0.16492167921014186
```

```

SPY_1: -0.08231192010382934
SPY_2: 0.14003125168721936
SPY_3: 0.16281501357531855
SPY_4: -0.22291642640919737

```

```

SPY_5: 0.032440993145706525
SPY_6: -0.3523263202214909
SPY_7: 0.3053225677321785
SPY_8: -0.08793690494411846
SPY_9: 0.40709688145805434
SPY_10: -0.1279959054390272
IEF_1: 0.033176543846550374
IEF_2: 0.01101114352522373
IEF_3: -0.004350955551808462
IEF_4: 0.03902488547750795
IEF_5: 0.050195221828548915
IEF_6: 0.007737613773306497
IEF_7: -0.018691900912255453
IEF_8: 0.09720196964699498
IEF_9: -0.2326844298836487
IEF_10: 0.10583673012105793

```

```

[[ 0 110]
 [ 0 145]]

```

Score: 0.5686274509803921

This is not a good model at all, as it is predicting **True** every single time. There is no real learning going on here, it just assumes that the SPY will always outperform IEF.

2.2 Question 2 (30spt)

2.2.1 Question 2.1

Using the same data, train/test split ratio, and consider the same classification problem as in Question 1.2. Create a feed-forward neural network with a single hidden layer (10 hidden nodes) densely connected to the inputs. You may choose any activation functions you wish.

```

[ ]: from tensorflow import keras
from tensorflow.keras.layers import Dense
import tensorflow as tf

tf.random.set_seed(personal)

model = keras.Sequential()
model.add(Dense(10, input_shape=(X_train.shape[1],), activation='tanh'))
model.add(Dense(1, activation='sigmoid'))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	210

```
-----
dense_1 (Dense)                (None, 1)                11
=====
```

```
Total params: 221
Trainable params: 221
Non-trainable params: 0
```

```
-----
2023-02-15 22:26:11.167900: I tensorflow/compiler/jit/xla_cpu_device.cc:41] Not
creating XLA devices, tf_xla_enable_xla_devices not set
2023-02-15 22:26:11.168061: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

2.2.2 Question 2.2

Train this neural network on the training data.

Evaluate the performance of this model by printing the confusion matrix and accuracy on the test data.

```
[ ]: model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=150, verbose=0, validation_split=0.
    ↪1)
```

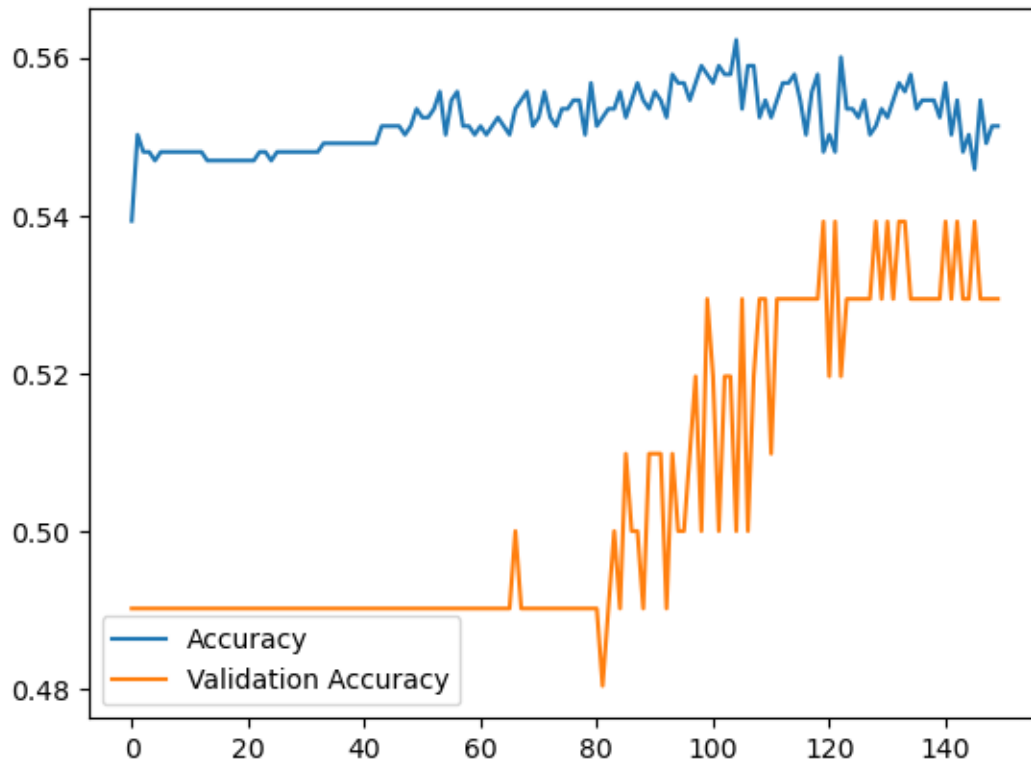
```
2023-02-15 22:26:11.279625: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR
optimization passes are enabled (registered 2)
```

```
[ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.show()
print(f"\n\nTrain Accuracy: {history.history['accuracy'][-1]}")
print(f"Validation Accuracy: {history.history['val_accuracy'][-1]}")

print(f"Test Accuracy: {model.evaluate(X_test, y_test, verbose=0)[1]}",
    ↪end='\n\n')

model_pred = (model.predict(X_test) > 0.5).astype(int)

print(confusion_matrix(y_test, model_pred))
```



Train Accuracy: 0.5513100624084473
 Validation Accuracy: 0.529411792755127
 Test Accuracy: 0.5764706134796143

```
[[ 23  87]
 [ 21 124]]
```

The accuracy is significantly better than the logistic regression, as it is actually making a decision based off the data. It is not always predicting True anymore and is getting good test accuracy.

2.2.3 Question 2.3

Using the same train/test split and consider the same classification problem as in Question 1.2. Train and test another feed-forward neural network of your own design.

```
[ ]: model12 = keras.Sequential()
model12.add(Dense(3, input_shape=(X_train.shape[1],), activation='selu'))
model12.add(Dense(3, activation='relu'))
model12.add(Dense(3, activation='relu'))
model12.add(Dense(3, activation='selu'))
model12.add(Dense(1, activation='sigmoid'))
```

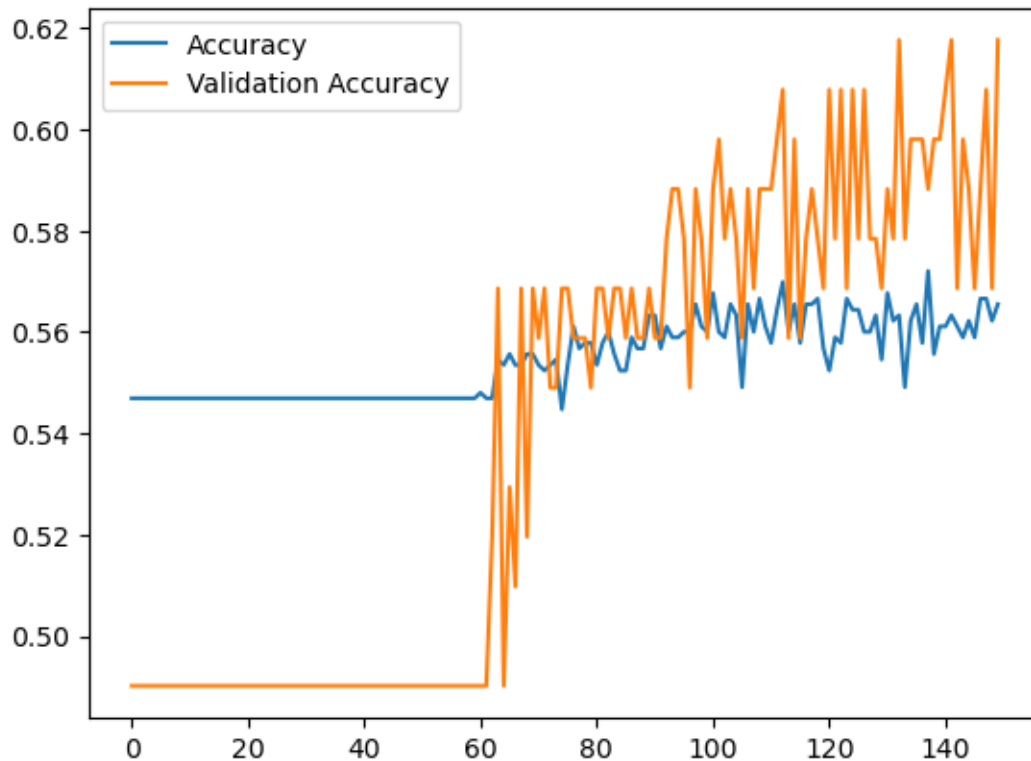
```
model2.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 3)	63
dense_3 (Dense)	(None, 3)	12
dense_4 (Dense)	(None, 3)	12
dense_5 (Dense)	(None, 3)	12
dense_6 (Dense)	(None, 1)	4

Total params: 103
Trainable params: 103
Non-trainable params: 0

```
[ ]: model2.compile(optimizer='adam', loss='binary_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history2 = model2.fit(X_train, y_train, epochs=150, verbose=0,  
    ↪ validation_split=0.1, batch_size=32)  
  
[ ]: plt.plot(history2.history['accuracy'], label='Accuracy')  
plt.plot(history2.history['val_accuracy'], label='Validation Accuracy')  
plt.legend()  
plt.show()  
print(f"\n\nTrain Accuracy: {history2.history['accuracy'][-1]}")  
print(f"Validation Accuracy: {history2.history['val_accuracy'][-1]}")  
  
print(f"Test Accuracy: {model2.evaluate(X_test, y_test, verbose=0)[1]}",  
    ↪ end='\n\n')  
  
model_pred = (model2.predict(X_test) > 0.5).astype(int)  
print(confusion_matrix(y_test, model_pred))
```



Train Accuracy: 0.5655021667480469
 Validation Accuracy: 0.6176470518112183
 Test Accuracy: 0.5254902243614197

```
[[43 67]
 [54 91]]
```

2.3 Question 3 (30pt)

2.3.1 Question 3.1

Using the same data, train/test split ratio, and consider the same classification problem as in Question 1.2. Create a time dilation neural network with a single convolutional layer (filter size of 10, kernel size of 2, dilation size of 1) densely connected to the inputs. You may choose any activation functions you wish.

Hint: The CNN can reference earlier lags on its own without feeding explicit memory inputs as was needed for the Question 2.

```
[ ]: model3 = keras.Sequential()
      model3.add(keras.layers.Conv1D(filters=10, kernel_size=2, dilation_rate=1,
      ↪ input_shape=(X_train.shape[1], 1)))
```



```

model3.add(keras.layers.Flatten())
model3.add(Dense(1, activation='sigmoid'))

model3.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 19, 10)	30
flatten (Flatten)	(None, 190)	0
dense_7 (Dense)	(None, 1)	191

Total params: 221
 Trainable params: 221
 Non-trainable params: 0

2.3.2 Question 3.2

Train this neural network on the training data. Evaluate the performance of this model by printing the confusion matrix and accuracy on the test data.

```

[ ]: model3.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
history3 = model3.fit(X_train.values.reshape(-1, X_train.shape[1], 1), y_train,
    ↪epochs=150, verbose=0, validation_split=0.1, batch_size=32)

```

```

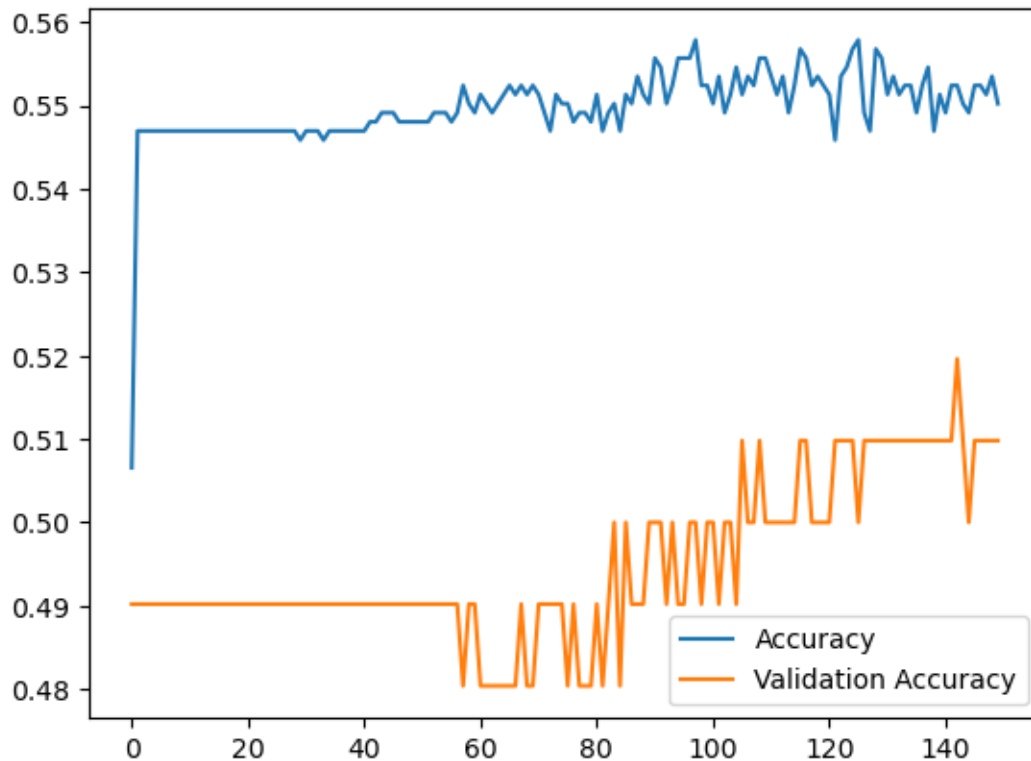
[ ]: plt.plot(history3.history['accuracy'], label='Accuracy')
plt.plot(history3.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.show()
print(f"\n\nTrain Accuracy: {history3.history['accuracy'][-1]}")
print(f"Validation Accuracy: {history3.history['val_accuracy'][-1]}")

print(f"Test Accuracy: {model3.evaluate(X_test.values.reshape(-1, X_test.
    ↪shape[1], 1), y_test, verbose=0)[1]}", end='\n\n')

model4_pred = (model3.predict(X_test.values.reshape(-1, X_test.shape[1], 1)) >
    ↪0.5).astype(int)

print(confusion_matrix(y_test, model4_pred))

```



Train Accuracy: 0.5502183437347412
 Validation Accuracy: 0.5098039507865906
 Test Accuracy: 0.5607843399047852

```
[[ 14  96]
 [ 16 129]]
```

2.3.3 Question 3.3

Using the same train/test split and consider the same classification problem as in Question 1.2. Train and test another convolutional neural network of your own design.

```
[ ]: model4 = keras.Sequential()
model4.add(keras.layers.Conv1D(filters=10, kernel_size=2, dilation_rate=1, u
    ↪ input_shape=(X_train.shape[1], 1)))
model4.add(keras.layers.Flatten())
model4.add(Dense(10, activation='relu'))
model4.add(Dense(1, activation='sigmoid'))

model4.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv1d_1 (Conv1D)	(None, 19, 10)	30
flatten_1 (Flatten)	(None, 190)	0
dense_8 (Dense)	(None, 10)	1910
dense_9 (Dense)	(None, 1)	11

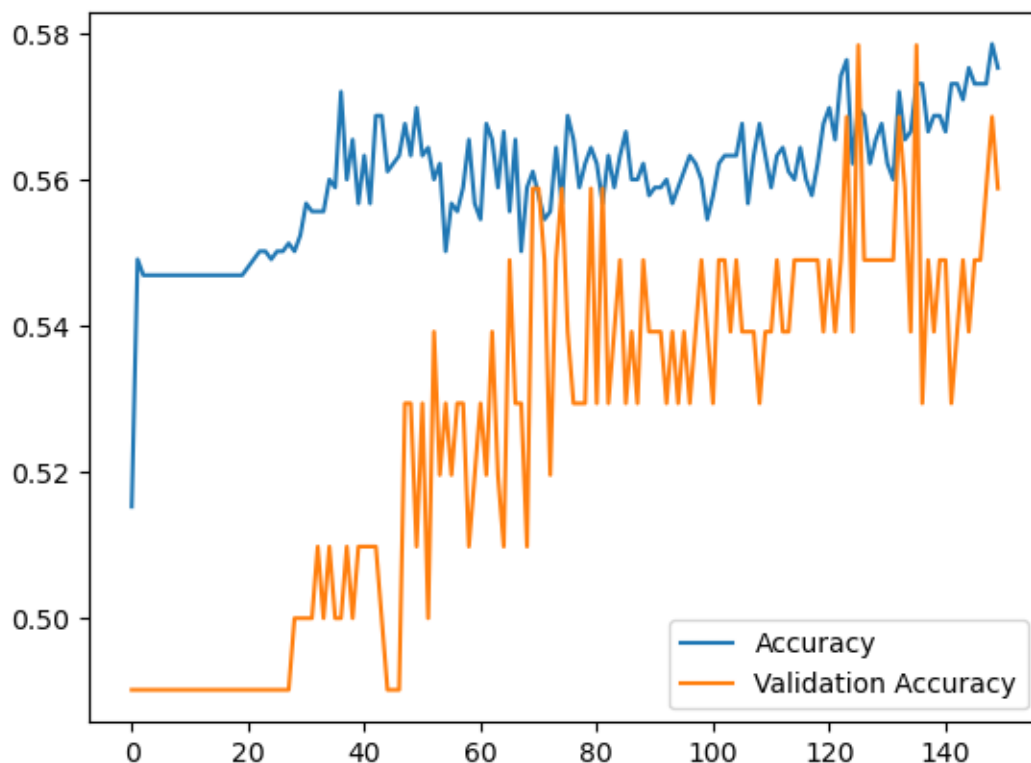
Total params: 1,951

Trainable params: 1,951

Non-trainable params: 0

```
[ ]: model4.compile(optimizer='adam', loss='binary_crossentropy',  
    ↪metrics=['accuracy'])  
  
history4 = model4.fit(X_train.values.reshape(-1, X_train.shape[1], 1), y_train,  
    ↪epochs=150, verbose=0, validation_split=0.1)
```

```
[ ]: plt.plot(history4.history['accuracy'], label='Accuracy')  
plt.plot(history4.history['val_accuracy'], label='Validation Accuracy')  
plt.legend()  
plt.show()  
print(f"\n\nTrain Accuracy: {history4.history['accuracy'][-1]}")  
print(f"Validation Accuracy: {history4.history['val_accuracy'][-1]}")  
  
print(f"Test Accuracy: {model4.evaluate(X_test.values.reshape(-1, X_test.  
    ↪shape[1], 1), y_test, verbose=0)[1]}", end='\n\n')  
  
model4_pred = (model4.predict(X_test.values.reshape(-1, X_test.shape[1], 1)) >  
    ↪0.5).astype(int)  
  
print(confusion_matrix(y_test, model4_pred))
```



Train Accuracy: 0.5753275156021118
 Validation Accuracy: 0.5588235259056091
 Test Accuracy: 0.5529412031173706

```
[[ 31  79]
 [ 35 110]]
```

2.4 Question 4 (30pt)

2.5 Question 4.1

Consider the same classification problem as in Question 1.2. Of the methods considered in this assignment, which would you recommend in practice? Explain briefly (1 paragraph) why you choose this fit.

The “base case” Logistic regression does not necessarily bring anything of value, as it is literally *always* predicting SPY to outperform IEF. This strategy would be to obviously buy and hold SPY for the end of time.

In terms of accuracy, the most simple feed-forward model whas the highest accuracy at over 57%, and it is not overfitting either. In fact, the test accuracy was even higher than the train.

The second CNN had roughly the same accuracy, but that could be because it is predicting True

at a significantly higher rate than the other, so that could be more so just luck as there are a lot more Trues than Falses in the test set. Therefore, I would choose the feed-forward model with 3 nodes at each hidden layer.

2.6 Question 4.2

Recreate your data set using data from January 1, 2022 through December 31, 2022. Using the method you would implement in practice, invest in the asset (SPY or IEF) depending on your predictions. Print the returns your portfolio would obtain from following this strategy. Comment on how this portfolio compares with the SPY and IEF returns and risks.

First, I will use buy-and-hold SPY as the base case to compare the strategy to.

```
[ ]: data = yf.download(tickers=tickers, start='2022-01-01', end='2022-12-31')['Adj_
    ↪Close']

ret = data.apply(lambda x: np.log(x / x.shift(1))) # log returns

lags = 10

for t in tickers:
    for i in range(1, lags+1):
        ret[f"{t}_{i}"] = ret[t].shift(i)

ret = ret.dropna()

print(ret.head(), end='\n\n')
print(f"{ret.shape = }")
```

```
[*****100%*****] 2 of 2 completed

      IEF      SPY      SPY_1      SPY_2      SPY_3      SPY_4 \
Date
2022-01-19  0.002414 -0.010438 -0.017868  0.000409 -0.013874  0.002700
2022-01-20  0.002141 -0.011130 -0.010438 -0.017868  0.000409 -0.013874
2022-01-21  0.005774 -0.019826 -0.011130 -0.010438 -0.017868  0.000409
2022-01-24 -0.000886  0.004238 -0.019826 -0.011130 -0.010438 -0.017868
2022-01-25 -0.001331 -0.012284  0.004238 -0.019826 -0.011130 -0.010438

      SPY_5      SPY_6      SPY_7      SPY_8 ...      IEF_1      IEF_2 \
Date
2022-01-19  0.009067 -0.001245 -0.003961 -0.000940 ... -0.007579 -0.006905
2022-01-20  0.002700  0.009067 -0.001245 -0.003961 ...  0.002414 -0.007579
2022-01-21 -0.013874  0.002700  0.009067 -0.001245 ...  0.002141  0.002414
2022-01-24  0.000409 -0.013874  0.002700  0.009067 ...  0.005774  0.002141
2022-01-25 -0.017868  0.000409 -0.013874  0.002700 ... -0.000886  0.005774

      IEF_3      IEF_4      IEF_5      IEF_6      IEF_7      IEF_8 \
```

Date						
2022-01-19	0.003181	0.000177	0.002393	-0.000266	-0.003277	-0.002472
2022-01-20	-0.006905	0.003181	0.000177	0.002393	-0.000266	-0.003277
2022-01-21	-0.007579	-0.006905	0.003181	0.000177	0.002393	-0.000266
2022-01-24	0.002414	-0.007579	-0.006905	0.003181	0.000177	0.002393
2022-01-25	0.002141	0.002414	-0.007579	-0.006905	0.003181	0.000177

	IEF_9	IEF_10
Date		
2022-01-19	-0.004224	-0.000527
2022-01-20	-0.002472	-0.004224
2022-01-21	-0.003277	-0.002472
2022-01-24	-0.000266	-0.003277
2022-01-25	0.002393	-0.000266

[5 rows x 22 columns]

ret.shape = (240, 22)

```
[ ]: # Calculate cumlulative returns of df['SPY']
spy_returns = np.zeros(len(ret['SPY'])+1)
spy_returns[0] = 1
for i in range(1, len(ret['SPY'])+1):
    spy_returns[i] = spy_returns[i-1] * (1 + ret['SPY'][i-1])

spy = spy_returns[-1] - 1
print(f"SPY Returns over the period: {spy}")
```

SPY Returns over the period: -0.17282694235256402

```
[ ]: test_data = ret.drop(columns=['SPY', 'IEF'])

predictions = (model.predict(test_data) > 0.5).astype(int).flatten()
predictions
```

```
[ ]: array([0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1,
          1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1,
          0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
          1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1,
          1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1,
          1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1,
          1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
          1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0,
          1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1,
          1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1])
```

```
[ ]: model_algo_daily_ret = np.zeros(len(ret['SPY']))

for i in range(len(model_algo_daily_ret)):
    if predictions[i] == 1:
        model_algo_daily_ret[i] = ret['SPY'][i]
    else:
        model_algo_daily_ret[i] = ret['IEF'][i]

model_algo_cum_ret = np.zeros(len(model_algo_daily_ret)+1)
model_algo_cum_ret[0] = 1
for i in range(1, len(model_algo_daily_ret)+1):
    model_algo_cum_ret[i] = model_algo_cum_ret[i-1] * (1 +
    ↪model_algo_daily_ret[i-1])

model_algo = model_algo_cum_ret[-1] - 1

print(f"Model Algo Returns over the period:          {model_algo}")

print(f"Difference compared to SPY (+ good, - bad): {model_algo - spy}",
    ↪end='\n\n')

print(f"SPY std:          {ret['SPY'].std() * np.sqrt(252)}")
print(f"Model Algo std: {np.std(model_algo_daily_ret) * np.sqrt(252)}",
    ↪end='\n\n')

print(f"SPY Sharpe:          {(ret['SPY'].mean() * 252) / (ret['SPY'].std() * np.
    ↪sqrt(252))}")
print(f"Model Algo Sharpe: {(model_algo_daily_ret.mean() * 252) / (np.
    ↪std(model_algo_daily_ret) * np.sqrt(252))}")
```

```
Model Algo Returns over the period:          -0.20946032861328945
Difference compared to SPY (+ good, - bad): -0.03663338626072543
```

```
SPY std:          0.24602843730869683
Model Algo std: 0.2261954165082829
```

```
SPY Sharpe:          -0.686808437379032
Model Algo Sharpe: -0.9771674907234906
```

You can see that the model underperformed the SPY by roughly 3%. It does come at a slightly lower volatility which makes sense as IEF has a lower vol. In terms of Sharpe Ratio, they are both negative, with the algo doing still worse. Therefore, it is not a good strategy to use and you are going to make more money by holding SPY over 2022.